

Understanding and Addressing Employee Attrition in the USA

Nazmus Saquib
nsaquib2
CSE 460
nsaquib2@buffalo.edu

Sakib Hassan
sakibhas
CSE 460
sakibhas@buffalo.com

Tawhidur Rahman
tawhidur
CSE 460
tawhidur@buffalo.edu

Overview of Phase 1

Problem Statement

Employee attrition in the USA has been on the rise over the past decade, posing significant challenges for organizations. The lack of understanding regarding the underlying factors contributing to attrition makes it difficult for businesses to implement effective retention strategies. Our project aims to analyze the reasons behind employee attrition and provide insights to help organizations mitigate this issue.

Background

The problem of employee attrition is significant as it impacts organizational productivity, morale, and ultimately, profitability. High turnover rates lead to increased recruitment and training costs, disrupt team dynamics, and can harm the company's reputation.

Contribution

By developing a database to analyze employee attrition, we aim to provide actionable insights for organizations to better understand and address this issue. Our contribution lies in offering a systematic approach to identify key factors driving attrition, thereby enabling organizations to implement targeted retention strategies.

Target Users:

1. Database Users

Human resource managers, organizational analysts, and business executives will use the database to gain insights into employee attrition trends, identify risk factors, and develop retention strategies.

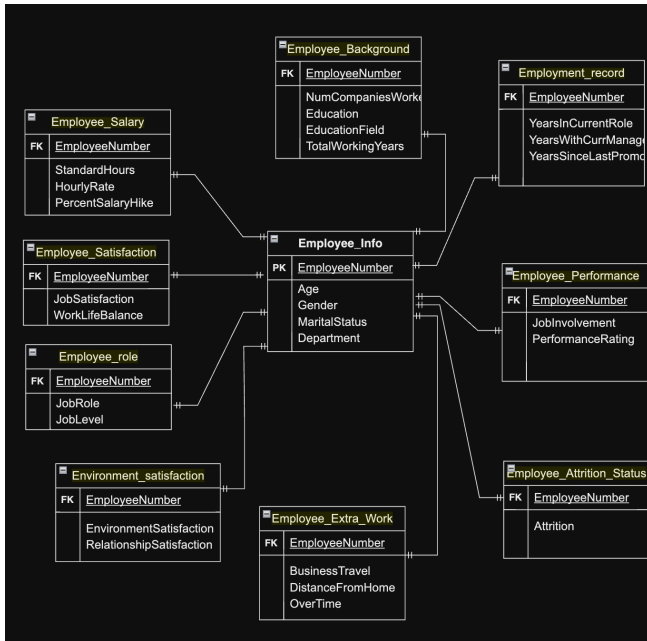
2. Database Administrators

The database administrators will be responsible for maintaining the database, ensuring data integrity, and optimizing query performance. They will also oversee data updates and security measures.

3. Real-Life Scenario

For instance, consider a multinational corporation with a high turnover rate in its call center operations. The HR department aims to understand the reasons behind this attrition and develop interventions to improve employee retention. Our database will serve as a valuable tool for HR managers to analyze employee data and make informed decisions.

E/R Diagram



The E/R diagram depicts the relationships between different tables in our database.

Here the employee number is a unique value for all the tables. We used it as a primary key for the employee_info table and as a foreign key for all other tables.

List of Relations and Attributes:

Employee_Info

Attributes: EmployeeNumber (Primary Key), Age, Gender, MaritalStatus, Department

Employee_Background:

Attributes: EmployeeNumber (Foreign Key), NumCompaniesWorked, Education, EducationField, TotalWorkingYears

Employee_Performance:

Attributes: EmployeeNumber (Foreign Key), JobInvolvement, PerformanceRating

Employee_Attrition_Status:

Attributes: EmployeeNumber (Foreign Key), Attrition

Employee_Extra_Work:

Attributes: EmployeeNumber (Foreign Key), BusinessTravel, DistanceFromHome, OverTime

Employee_Satisfaction:

Attributes: EmployeeNumber (Foreign Key), JobSatisfaction, WorkLifeBalance,

Employee_Salary:

Attributes: EmployeeNumber (Foreign Key), StandardHours, HourlyRate, PercentSalaryHike

Employee_role:

Attributes: EmployeeNumber (Foreign Key), JobRole, JobLevel

Employee_record:

Attributes: EmployeeNumber (Foreign Key), YearsInCurrentRole, YearsWithCurrManager, YearsSinceLastPromotion

Environment_satisfaction:

Attributes: EmployeeNumber, (Foreign Key) EnvironmentSatisfaction, RelationshipSatisfaction

Schema:

Employee_Info(EmployeeNumber (Primary Key) varchar, Age int, Gender varchar, MaritalStatus varchar, Department varchar)

Employee_Background(EmployeeNumber (Foreign Key) varchar,

NumCompaniesWorked int, Education
varchar, EducationField varchar,
TotalWorkingYears int)

Employee_Performance(EmployeeNumber
(Foreign Key) varchar, JobInvolvement
varchar, PerformanceRating int)

Employee_Attrition_Status(EmployeeNumber
(Foreign Key) varchar, Attrition varchar)

Employee_Extra_Work(EmployeeNumber
(Foreign Key) varchar, BusinessTravel int,
DistanceFromHome int, OverTime int)

Employee_Satisfaction(EmployeeNumber
(Foreign Key) varchar, JobSatisfaction int,
WorkLifeBalance int)

Employee_Salary(EmployeeNumber (Foreign
Key) varchar, StandardHours int, HourlyRate
int, PercentSalaryHike int)

Employee_role(EmployeeNumber (Foreign
Key) varchar, JobRole varchar, JobLevel int)

Employee_record(EmployeeNumber (Foreign
Key) varchar, YearsInCurrentRole int,
YearsWithCurrManager int,
YearsSinceLastPromotion int)

Environment_satisfaction(EmployeeNumber
(Foreign Key) varchar,
EnvironmentSatisfaction int,
RelationshipSatisfaction int)

BCNF Transformation: We ensured that all
relations are in Boyce-Codd Normal Form
(BCNF) by checking any functional
dependencies or redundancies. This involved

decomposing tables or modifying their
structure to eliminate anomalies and maintain
data integrity.

Data Acquisition and Preparation

We acquired a large dataset containing relevant
information about employee attributes, such as
demographics, job satisfaction, performance
ratings, and work-life balance. This dataset is
obtained from real-world sources such as HR
databases or simulated using data generation
programs. We ensured that the dataset aligns
with our database schema and contains enough
diversity to support various types of SQL
queries.

Data Acquisition: Obtained a large dataset
containing information on employee attributes
such as age, gender, education level, job role,
performance ratings, etc. This dataset reflected
real-world scenarios and was sufficiently
diverse to capture different factors influencing
employee attrition.

Data Preparation: Ensured that the acquired
dataset was in a suitable format for loading
into our database. This involved writing scripts
to transform the data into a compatible form,
cleaning the data to remove any
inconsistencies or errors, and structuring the
data to fit our database schema.

Data Loading: Loaded the dataset into our
database, ensuring that each table corresponds
to a specific entity or attribute of interest.
Verified that the data is accurately represented
in the database and that relationships between
tables are properly established.

Data Validation: Performed validation checks to ensure the integrity and quality of the data. This involved running queries to verify data consistency, identifying any anomalies or discrepancies, and resolving any issues that arise.

Creating Links: Ensured that the dataset contained enough interconnections between rows of different tables to support a variety of advanced SQL queries.

Ensuring Boyce-Codd Normal Form (BCNF) Compliance

To ensure that our database design adheres to BCNF, we analyzed the functional dependencies within each relation and decomposed tables as necessary to eliminate any redundancy or dependency violations.

Dependency Analysis: We identified the functional dependencies within each relation in our database schema.

BCNF Decomposition: No relation is found to violate BCNF, so we did not have to decompose the relation into smaller relations to eliminate redundancy and dependency violations.

In our initial schema, there is no multi-valued attribute, no partial dependency, and no transitive dependency, where EmployeeNumber is the superkey. So, all the relations are in BCNF.

Why the database is in BCNF:

Primary Key Uniqueness: In BCNF, each table must have a primary key, and all non-key

attributes must be non-transitively dependent (i.e., directly dependent) on the primary key. In our diagram, EmployeeNumber is indicated as the primary key (PK) for all tables. This suggests that each table's attributes are functionally dependent on EmployeeNumber, which is essential for BCNF compliance.

Functional Dependencies: The relationships (arrows) between tables primarily reflect foreign key (FK) constraints where EmployeeNumber in one table links to EmployeeNumber in another, signifying that all attribute dependencies on EmployeeNumber are direct and not through some other attribute. This direct dependency aligns with BCNF's requirement that every determinant in a table is a candidate key.

Normalization Check:

Employee_Info: Attributes like Age, Gender, MaritalStatus, and Department directly depend on EmployeeNumber. There are no attributes that functionally determine EmployeeNumber, nor are there transitive dependencies.

Employee_Performance and others: Similarly, in tables like Employee_Performance, Employee_Salary, Employee_Satisfaction, and more, each non-key attribute is functionally dependent only on EmployeeNumber. This eliminates the possibility of transitive dependencies where a non-key attribute determines another non-key attribute.

No Redundancy or Update Anomalies: Each piece of information is stored only once. For instance, employee age is stored only in the Employee_Info table and not repeated

elsewhere, which prevents update anomalies and maintains data integrity.

The ER diagram, therefore, exhibits a well-structured relational schema in which all non-key attributes in every table are directly dependent on the primary keys, with no transitive dependencies. This is the core requirement for a database to be in BCNF. Such a design helps in maintaining data consistency and integrity across the database, which is crucial for effective data management and retrieval.

Phase 2

1. Handling Larger Dataset and Indexing Concepts

Challenges Encountered

One challenge we faced when we dealt with our large dataset was slower query performance due to the volume of data being processed. Additionally, loading large datasets into the database took considerable time and required optimization to ensure efficient data ingestion.

Solution through indexing

To address performance issues, we adopted indexing concepts to optimize query execution. Indexes were created on columns frequently used in WHERE clauses or JOIN conditions to speed up data retrieval.

We identified key columns in our dataset that were frequently used in queries, such as EmployeeNumber and Age. Indexes were created on these columns to improve query

performance. An example of our indexing query is below.

```
Query Query History
1 Update employee_background set education = '2' where employeeenumber = '34';
2
3
4
Data Output Messages Notifications
UPDATE 0
Query returned successfully in 163 msec.
```

Here we can see that the query took 163 mili sec. Then we applied indexing on the EmployeeNumber column of the employee_background table.

```
Query Query History
1
2
3 CREATE INDEX idx_emp_bg
4 ON employee_background (Employeeenumber);
Data Output Messages Notifications
CREATE INDEX
Query returned successfully in 286 msec.
```

Below is the result after indexing.

```
Query Query History
1 Update employee_background set education = '2' where employeeenumber = '34';
2
3 CREATE INDEX idx_emp_bg
4 ON employee_background (Employeeenumber);
Data Output Messages Notifications
UPDATE 0
Query returned successfully in 44 msec.
```

We can see that the same query that took 163 mili sec before now takes only 44 mili sec after indexing. It means the query was optimized by 70%.

Regular monitoring of query execution times helped us evaluate the effectiveness of indexing strategies. Adjustments were made as needed to further optimize performance. By adopting these indexing concepts, we were

able to enhance the efficiency of our database queries and mitigate the challenges posed by handling larger datasets.

2. Testing database with 10 queries

Inserting Data (2 queries):

2.1. Inserted a new employee record into the employee_info table.

QueryQuery History

1select * from employee_info

2

3INSERT INTO employee_info (employeenumber, age, gender, maritalstatus, department

4VALUES ('7776', 27, 'Female', 'Single', 'Marketing');

5

6

7

Data OutputMessagesNotifications

INSERT 0 1

Query returned successfully in 53 msec.

2.2. Inserted a record into the employee_extra_work table for a specific employee.

QueryQuery History

1select * from employee_extra_work where employeenumber = '1';

2

3INSERT INTO employee_extra_work (employeenumber, businesstravel, distancefromhome, overtime)

4VALUES ('7776', 'Travel_Rarely', '2', 'Yes');

5

6

7

Data OutputMessagesNotifications

INSERT 0 1

Query returned successfully in 40 msec.

Delete (2 queries):

2.3. Deleted an employee record from the employee_extra_work table.

QueryQuery History

1select * from employee_extra_work;

2

3DELETE FROM employee_extra_work WHERE EmployeeNumber = '7776';

4

5

6

7

Data OutputMessagesNotifications

DELETE 1

Query returned successfully in 89 msec.

2.4. Deleted employee records whose medical field are medical from the employee_background table.

QueryQuery History

1select * from employee_background;

2

3DELETE FROM employee_background WHERE educationfield = 'Medical';

4

5

6

7

Data OutputMessagesNotifications

DELETE 464

Query returned successfully in 56 msec.

Update (2 queries):

2.5. Updated education to “1” for a specific employee in the employee_background table.

QueryQuery History

1select * from employee_background where EmployeeNumber = '23';

2

3update employee_background set education = '1' where EmployeeNumber = '23';

4

5

6

7

Data OutputMessagesNotifications

	employeenumber character varying (30)	numcompaniesworked character varying (30)	education character varying (30)	educationfield character varying (30)	totalworkingyears character varying (30)
1	23	2	4	Life Sciences	31

Query	Query History
1	<code>select * from employee_background where EmployeeNumber = '23';</code>
2	
3	<code>update employee_background set education = '1' where EmployeeNumber = '23';</code>
4	
5	
6	
7	

Data Output	Messages	Notifications
UPDATE 1		
Query returned successfully in 44 msec.		

2.6. Updated employee_role to “IT executive” for a specific employee in the employee_role table.

Query	Query History
1	<code>select * from employee_role where EmployeeNumber = '445';</code>
2	
3	<code>update employee_role set jobrole = 'IT Executive' where EmployeeNumber = '445';</code>
4	
5	
6	
7	

Data Output	Messages	Notifications												
<table> <tr> <th>employeenumber</th><th>jobrole</th><th>joblevel</th></tr> <tr> <td>character varying (30)</td><td>character varying (30)</td><td>character varying (30)</td></tr> <tr> <td>1</td><td>445</td><td>Sales Executive</td></tr> <tr> <td></td><td></td><td>2</td></tr> </table>	employeenumber	jobrole	joblevel	character varying (30)	character varying (30)	character varying (30)	1	445	Sales Executive			2		
employeenumber	jobrole	joblevel												
character varying (30)	character varying (30)	character varying (30)												
1	445	Sales Executive												
		2												

Query	Query History
1	<code>select * from employee_role where EmployeeNumber = '445';</code>
2	
3	<code>update employee_role set jobrole = 'IT executive' where EmployeeNumber = '445';</code>
4	
5	
6	
7	

Data Output	Messages	Notifications
UPDATE 1		
Query returned successfully in 83 msec.		

Select (4 queries):

2.7. Joined Employee_info and Employee_Performance table on EmployeeNumber and selected age, department, performancerating, and employeenumber for all the employees.

Query

Query History

1

2

3

4

```
SELECT e.EmployeeNumber, e.Age, e.Department, ep.PerformanceRating
FROM Employee_info e
JOIN Employee_Performance ep ON e.EmployeeNumber = ep.EmployeeNumber;
```

Data Output

Messages

Notifications

	employeenumber character varying (30)	age character varying (30)	department character varying (30)	performancerating character varying (30)
1	1	41	Sales	3
2	2	49	Research & Development	4
3	4	37	Research & Development	3
4	5	33	Research & Development	3
5	7	27	Research & Development	3
6	8	32	Research & Development	3

2.8. Calculated the average age for each department (group by) in the Employee_info table, converting the Age column to a numeric type before averaging to handle potential string representations of age values.

Query

Query History

1

2

3

4

SELECT Department, AVG(Age::numeric) AS AverageAge

FROM Employee_info

GROUP BY Department;

Data Output

Messages

Notifications

department

character varying (30)

averageage

numeric

1

Marketing

27.0000000000000000

2

Human Resources

37.8095238095238095

3

Research & Development

37.0426638917793965

4

Sales

36.5426008968609865

2.9. Retrieved all columns (*) from the Employee_info table and ordered the results in descending (DESC) order based on the Age column.

Query

Query History

1

2

3

4

SELECT * FROM Employee_info ORDER BY Age DESC;

Data Output

Messages

Notifications

	employeenumber [PK] character varying (30)	age character varying (30)	gender character varying (30)	maritalstatus character varying (30)	departmen character varying (30)
1	549	60	Female	Married	Research
2	732	60	Male	Single	Sales
3	1697	60	Male	Divorced	Research
4	1233	60	Male	Divorced	Sales
5	573	60	Female	Married	Sales
6	10	59	Female	Married	Research
7	321	59	Male	Married	Human Resources
8	81	59	Female	Single	Sales
9	140	59	Female	Married	Human Resources

2.10. Selected the maximum value of the PerformanceRating column from the Employee_Performance table and renamed the result as HighestPerformanceRating. It calculates the highest performance rating recorded in the table.

Query Query History	
1	SELECT MAX(PerformanceRating) AS HighestPerformanceRating FROM Employee_Performance;
2	
3	
4	
Data Output Messages Notifications	
<div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> <div></div> </div>	
	highest_performance_rating text
1	4

3. Query execution analysis

We analyzed the explanation generated by PostgreSQL for each query to identify potential bottlenecks and areas for optimization.

Problematic Query 1:

Query Query History	
1	SELECT MAX(PerformanceRating) AS HighestPerformanceRating FROM Employee_Performance;
2	
3	
4	
Data Output Messages Explain X Notifications	
Graphical Analysis Statistics	
#	Node
1.	→ Aggregate
2.	→ Seq Scan on employee_performance as employee_performance

Using the explain tool we found out that one way to optimize this query is to create an index on the PerformanceRating column in the Employee_Performance table. This can improve the performance of the query by allowing PostgreSQL to quickly find the maximum value without scanning the entire table.

Problematic query 2:

Query Query History	
1	SELECT e.EmployeeNumber, e.Age, e.Department, ep.PerformanceRating
2	FROM Employee_info e
3	JOIN Employee_Performance ep ON e.EmployeeNumber = ep.EmployeeNumber;
4	
Data Output Messages Explain X Notifications	
Graphical Analysis Statistics	
#	Node
1.	→ Hash Inner Join Hash Cond: ((ep.employee_number)::text = (e.employee_number)::text)
2.	→ Seq Scan on employee_performance as ep
3.	→ Hash
4.	→ Seq Scan on employee_info as e

One way to optimize the query is to use Integer Keys and ensure that EmployeeNumber in both tables is of the same integer type for faster join performance. Moreover, we can create indexes on EmployeeNumber in both Employee_info and Employee_Performance tables to speed up the join operation. Thirdly, we can avoid converting EmployeeNumber to text for the join condition. We need to ensure that both

columns are of the same data type. Lastly, we can also consider using a WHERE clause to filter the data before joining to reduce the number of rows processed.

Problematic Query 3:

Query		Query History
1	SELECT Department, AVG(Age::numeric) AS AverageAge	
2	FROM Employee_info	
3	GROUP BY Department;	
4		
Data Output		Messages
		Explain X
		Notifications
Graphical		Analysis
		Statistics
#	Node	
1.	→ Aggregate	
2.	→ Seq Scan on employee_info as employee_info	

With the help of “explain” tool and referring to the book, we came to know that we can optimize this query by ensuring that the ‘Age’ column is stored as a numerical data type (INTEGER, FLOAT, etc.) instead of character varying to avoid the need for casting. Furthermore, we could create an index on the Department column to speed up the grouping process.

Conclusion

By analyzing the execution plans and proposing optimizations, we can address potential performance issues in the identified queries and improve overall database efficiency.

Contributions

33% Each member of the team.